

Playing card detection

PlayCDC

Object recognition and image understanding lecture

Frank Gabel and Daniel Gonzalez

15.07.2018

1 Abstract

With the capabilities of upcoming small video capturing devices in, for example, smart contact lenses with built-in cameras, whole new ways of cheating in cardgames emerge. In order to help facilitate these cheating endeavours, we implement an algorithm that detects the suits and ranks of playing cards in the field of view of a camera using the latest iteration of the YOLO object recognition algorithm.



Figure 1: Stock photos of an envisioned camera-equipped contact lense (left) and a mysterious black-jack player (right)

2 Introduction

A big topic in computer vision is the search of methods that are, given a query image, capable of answering questions like: What is present in an image? Is there a particular object in it? Where exactly in the image is this object located? Is it possible to semantically segment objects of interest in the image? Object detection deals with detecting instances of semantic objects of a certain class (such as humans, buildings, or cars) in digital images and videos. Typically, object detection deals with two sub-tasks: **object localization** using bounding boxes and **multiclass object classification** within said bounding boxes. Sometimes, a third sub-task of **semantic object segmentation** is performed, i.e. the process of labelling objects on pixel level.

In this project, we use the YOLOv3 object detection algorithm (CITE) for object detection (without additional segmentation) in order to tell apart playing cards of a standard 52-part deck.

3 Dataset

DANIEL THAT'S YOUR PLAYGROUND, PLEASE PUT A COUPLE OF PHOTOS :))) As there was no suitable dataset of cards to use, we decided to create it by our own. We did this in basically three main steps to obtain at the end (**Nr. images**) images, as well as the necessary Bounding

Box information required for training our YOLO NN. First we had to create the data doing photos and arrange it in such a way that we can easily use it around our workflow. The next, we had to detect the convex hulls of the card numbers and suits and the last perform rotations/re scaling/... on the images to generate new data.

3.1 dataset synthesis pipeline

Data preparation

We decided to work with a deck of 52 cards, where each card contain two times the card logo. We took for each card 2 different photos (**TODO: check this there were not really 52**) and after this we extract the cards of the photos and re-scaled them to 600x900 pixels. We choose this resolution arbitrary, as we though that it was a nice size to work with. For this part, we used the selection/rotation/cropping/re-scale -tools provided by GIMP¹. Concluding this, we proceed by detecting the convex hulls of the cards numbers and suits. We detect the convex hulls using SciPy libraries² in a semi-automatic way, verifying for each card manually that we got the desired result [SEE FIGURE]. At the last step, we saved the convex hulls as an NumPy array (**cite??**)

Data generation

The goal in this step was to generate a big amount of new data for each card destined for training or NN. We wanted to perform image transformations, like translations, rotations and re scaling, as well as blurring and sharpening to generate new data. We decided to use the imgaug python library, because it provided a nice way to keep track of the convex hulls after doing transformations.

4 Related work - The object detection landscape

The task of object detection in images encompasses both the "simple" form of localizing and subsequently classifying objects as well as the more "advanced" form of additionally segmenting objects pixel-wise. For these task, different deep learning architectures have emerged, of which we will present the most important ones in the following.

In 2015, **Faster Region-based Convolutional Network (Faster R-CNN)** (CITE) has come up as an enhancement of the existing R-CNN and Fast R-CNN methods that are based on both region proposal networks to find candidate bounding boxes and detection networks to perform classification. Faster R-CNN extend this by introducing weight sharing of the convolutional features between region proposal network and detection network, facilitating nearly cost-free region proposals.

You Only Look Once (YOLO) Single-Shot Detector (SSD)

Focal Loss for Dense Object Detection (RetinaNet) Mask Region-based Convolutional Network (Mask R-CNN) The Mask-RCNN consitutes the current state-of-the-art in semantic object segmentation. The basic idea is to..... As this extends the scope of our project, we will not dive deeper and rather refer interested readers to the original research paper (CITE) or ...

5 Methods

The YOLO approach to object detection

The YOLO model's novel motivation is that it re-frames object detection as a single regression problem, directly from image pixels to bounding box coordinates and class probabilities. This means that the YOLO model only "looks once" at an image for object detection. It works by dividing an image into $S \times S$ grid. Each grid cell predicts B bounding boxes and a confidence score for each box. The confidence scores show how confident the model is that there is an object in that bounding box. This is formally defined as:

¹GIMP 2.8.22 - GNU Image Manipulation Program

²SciPy: Open Source Scientific Tools for Python




Figure 2: tiny YOLO-v3 feature extractor

techniques

Non-max suppression Non max suppression removes the low probability bounding boxes which are very close to a high probability bounding boxes.

tiny YOLO-v3

The particular architecture we used for training on our dataset is **tiny YOLOv3** which essentially is a smaller version of YOLO for constrained environments.

5.1 Webcam deployment

DANIEL

5.2 Evaluation

DANIEL - WRITE SOMETHING ABOUT MAP, YOU CAN PUT FORMULAS IN HERE LIKE A BOSS, BUILD UPON WHAT I WROTEIN THE POSTER MAYBE

6 Results

Our object detection endeavours using a tiny YOLOv3 detector resulted in diverse results depending on the complexity and peculiarities of the synthesized dataset in use - the main results are summarize in REF TABLE 1 and REF FIGURE X. We generally achieved the best results for the datasets without textures and In total, training always converged rather quickly. Again, datasets without complex textures led to better (i.e. faster) convergence which intuitively makes sense as the algorithm does not need to spend time ruling out complex dependencies that might be learned from unimportant pixels or regions outside the card.

Dataset situation		Precision	Recall	mAP
Dataset	Description			
1 - Simple	Paste cards on simple canvases	2		
	<i>random rotations, brightness, blurring</i>	13		
2 - Medium	Paste randomly scaled cards on simple canvases	2		
	<i>random rotations, brightness, blurring</i>	13		
3 - Elaborate	Paste randomly scaled cards on textures	2		
	<i>random rotations, brightness, blurring</i>	13		
4 - Hardest	Paste randomly scaled cards on textures	2		
	<i>random rotations, brightness, blurring</i>	13		

Table 1: Your caption here

Extensions and tricks used in training

7 Discussion

FRANK and DANIEL

8 Conclusion

FRANK and DANIEL